
Sparse Feature Learning

Philipp Koehn

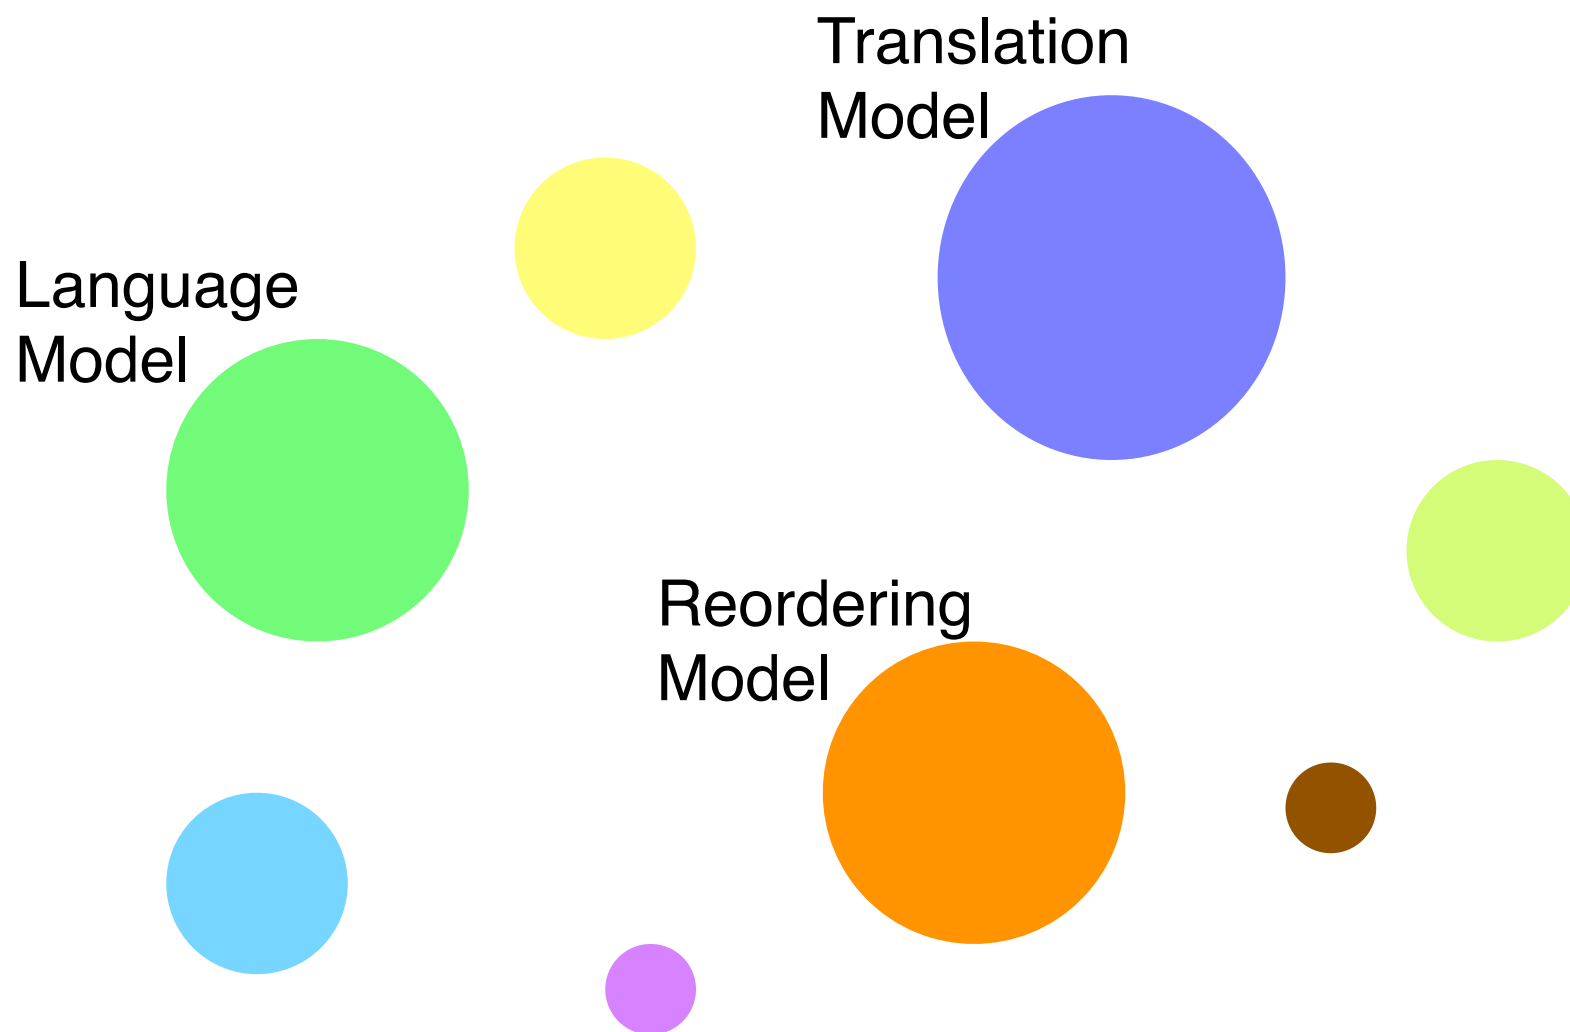
3 March 2015



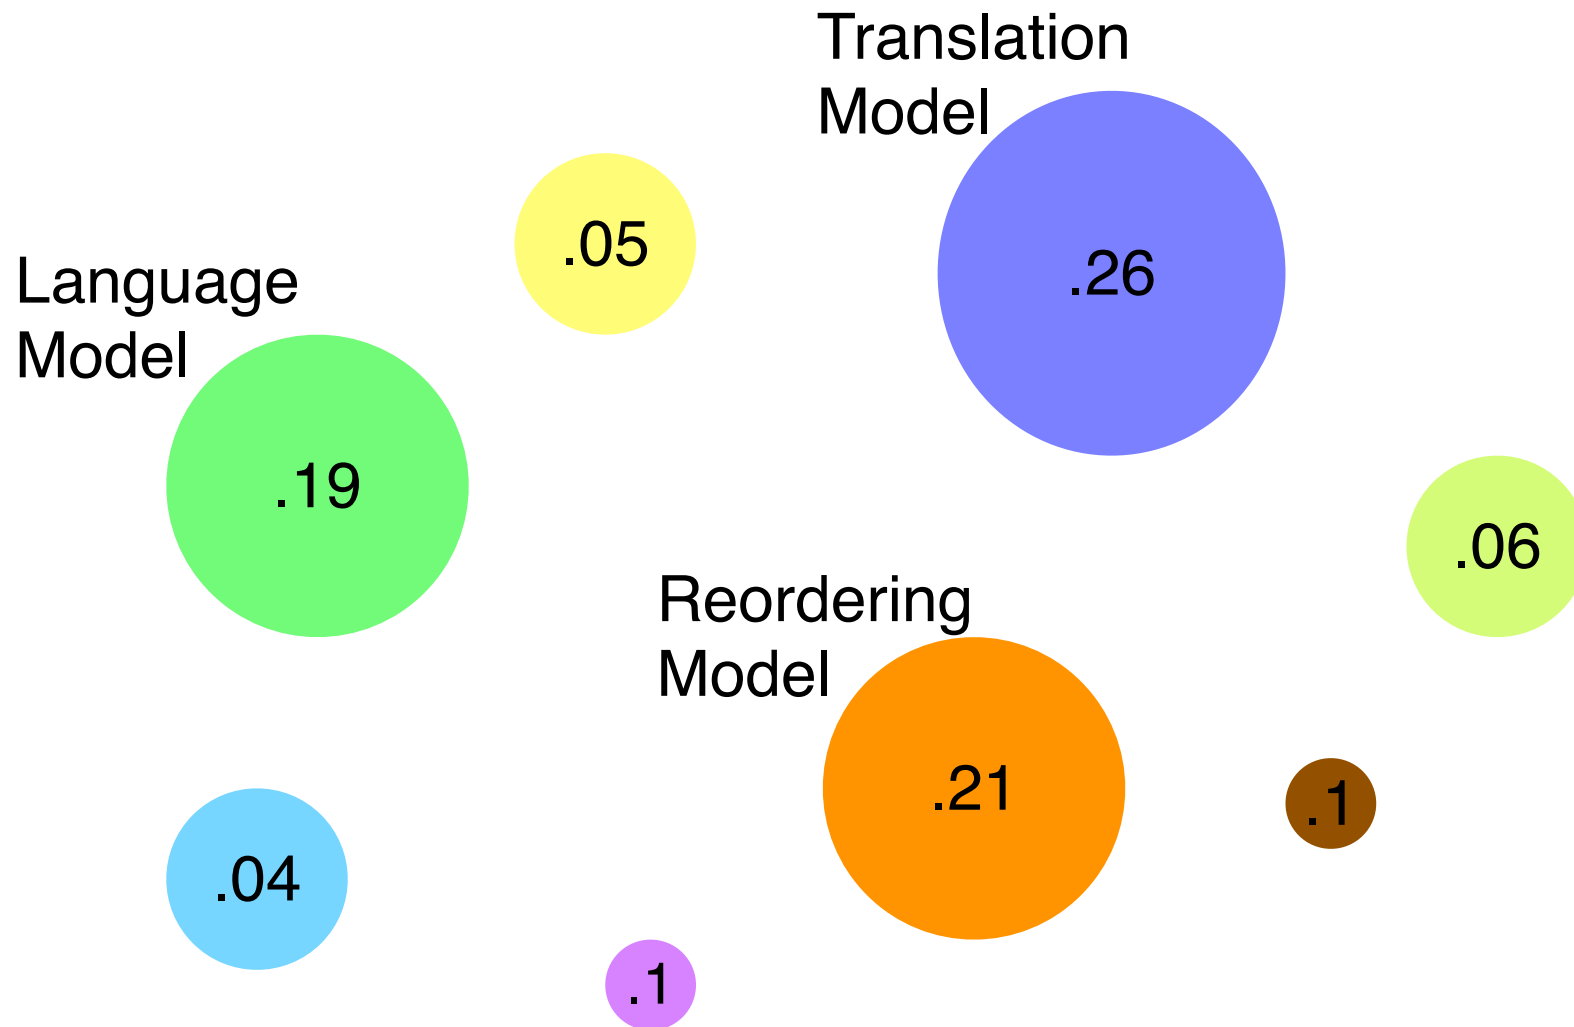
Multiple Component Models



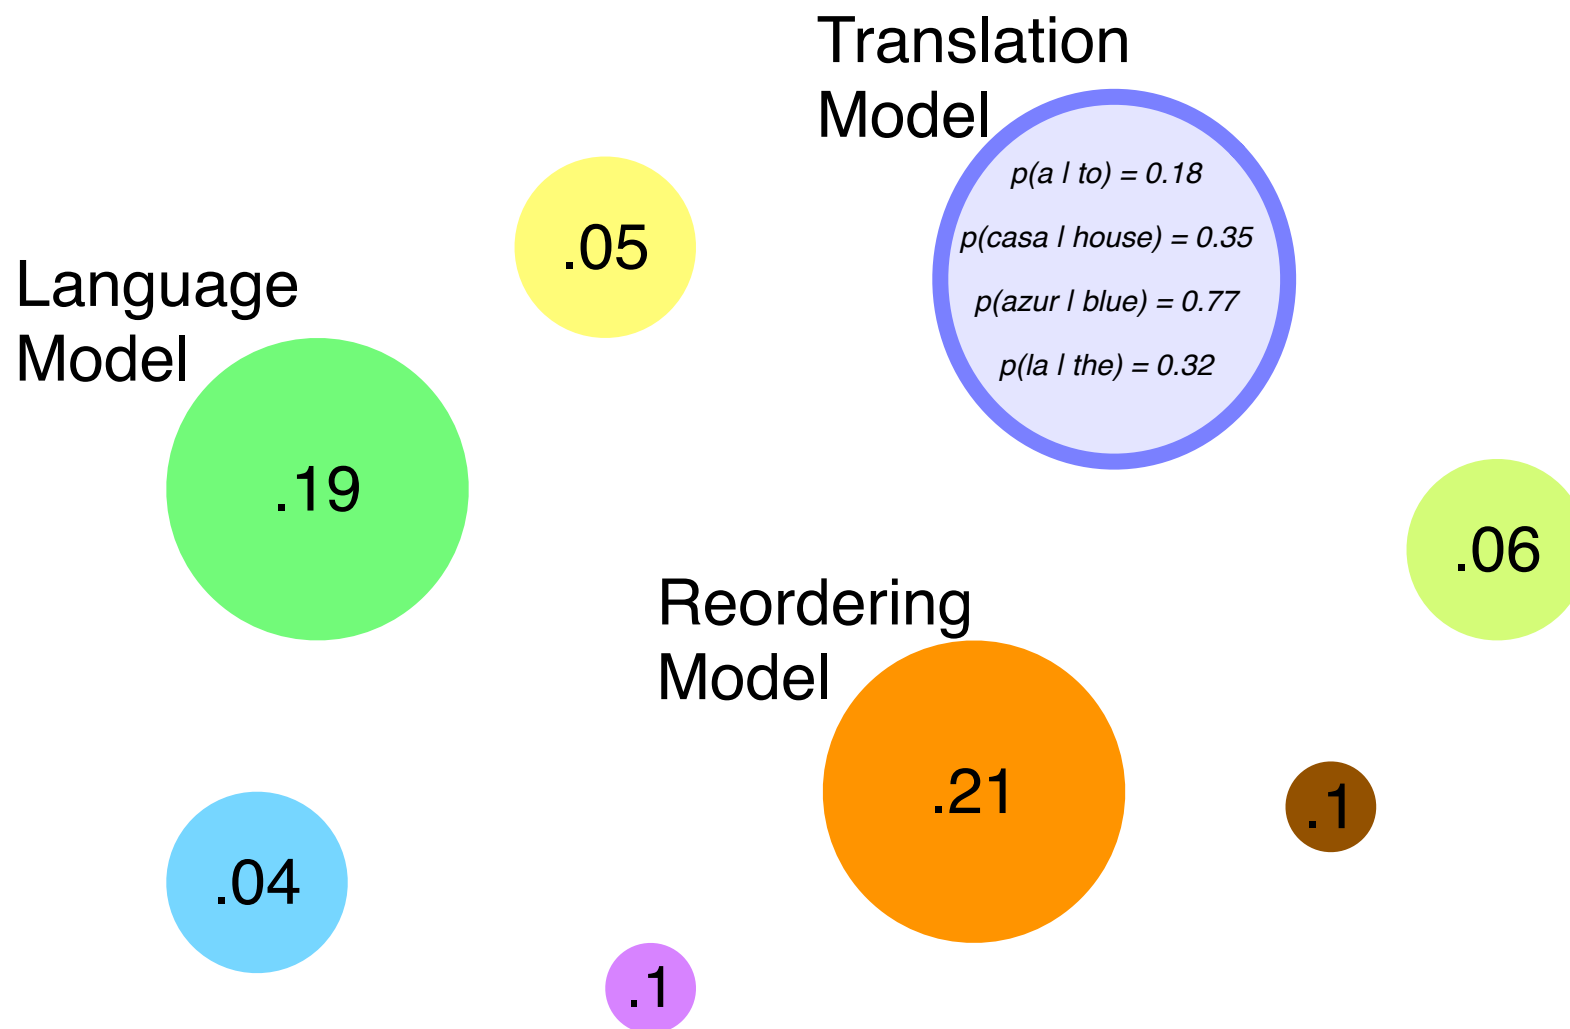
1



Component Weights




Even More Numbers Inside





Grand Vision





- There are millions of parameters
 - each phrase translation score
 - each language model n-gram
 - etc.
- Can we train them all discriminatively?
- This implies optimization over the entire training corpus











Strategy and Core Problems

- Process each sentence pair in the training corpus
- Optimize parameters towards producing the reference translation
- Reference translation may not be producible by model
 - optimize towards most similar translation
 - or, only process sentence pair partially
- Avoid overfitting
- Large corpora require efficient learning methods

Sentence Level vs. Corpus Level Error Metric¹¹



- Optimizing BLEU requires optimizing over the entire training corpus

$$\text{score}(\{\mathbf{e}_i^* = \operatorname{argmax}_{\mathbf{e}_i} \sum_j h_j(\mathbf{e}_i, \mathbf{f}_i) \lambda_i\}, \{\mathbf{e}_i^{\text{ref}}\})$$

- Life would be easier, if we could sum over sentence level scores

$$\sum_i \text{score}(\operatorname{argmax}_{\mathbf{e}_i} \sum_j h_j(\mathbf{e}_i, \mathbf{f}_i) \lambda_i, \mathbf{e}_i^{\text{ref}})$$

- For instance, BLEU+1

features

Core Rule Properties

- Frequency of phrase (binned)
- Length of phrase
 - number of source words
 - number of target words
 - number of source and target words
- Unaligned / added (content) words in phrase pair
- Reordering within phrase pair

- $\text{lex}(e)$ fires when an output word e is generated
- $\text{lex}(f, e)$ fires when an output word e is generated aligned to a input word f
- $\text{lex}(\text{NULL}, e)$ fires when an output word e is generated unaligned
- $\text{lex}(f, \text{NULL})$ fires when an input word e is dropped
- Could also be defined on part of speech tags or word classes

Lexicalized Reordering Features

- Replacement of lexicalized reordering model
- Features differ by
 - lexicalized by first or last word of phrase (source or target)
 - word representation replaced by word class
 - orientation type

- Indicator feature that the rule occurs in one specific domain
- Probability that the rule belongs to one specific domain
- Domain-specific lexical translation probabilities

- If we have syntactic parse trees, many more features
 - number of nodes of a particular kind
 - matching of source and target constituents
 - reordering within syntactic constituents
- Parse trees are a by-product of syntax-based models
- More on that in future lectures

Every Number in Model

- Phrase pair indicator feature
- Target n-gram feature
- Phrase pair orientation feature


perceptron algorithm

Optimizing Linear Model

- We consider each sentence pair $(\mathbf{e}_i, \mathbf{f}_i)$ and its alignment \mathbf{a}_i
- To simplify notation, we define derivation $\mathbf{d}_i = (\mathbf{e}_i, \mathbf{f}_i, \mathbf{a}_i)$
- Model score is weighted linear combination of feature values h_j and weights λ_j

$$\text{score}(\lambda, \mathbf{d}_i) = \sum_j \lambda_j h_j(\mathbf{d}_i)$$

- Such models are also known as single layer perceptrons



- Besides the reference derivation \mathbf{d}_i and its score

$$\text{score}(\lambda, \mathbf{d}_i) = \sum_j \lambda_j h_j(\mathbf{d}_i)$$

- We also have the model best translation

$$\mathbf{d}_i^{\text{best}} = \operatorname{argmax}_{\mathbf{d}} \text{score}(\lambda_i, \mathbf{d}_i) = \operatorname{argmax}_{\mathbf{d}} \sum_j \lambda_j h_j(\mathbf{d})$$

- ... and its model score

$$\text{score}(\lambda, \mathbf{d}_i^{\text{best}}) = \sum_j \lambda_j h_j(\mathbf{d}_i^{\text{best}})$$

- We can view the error in our model as a function of its parameters λ

$$\text{error}(\lambda, \mathbf{d}_i) = \text{score}(\lambda, \mathbf{d}_i^{\text{best}}) - \text{score}(\lambda, \mathbf{d}_i)$$

Stochastic Gradient Descent

- We want to minimize the error

$$\text{error}(\lambda, \mathbf{d}_i) = \text{score}(\lambda, \mathbf{d}_i^{\text{best}}) - \text{score}(\lambda, \mathbf{d}_i)$$

- In stochastic gradient descent, we follow direction of gradient

$$\frac{d}{d \lambda} \text{error}(\lambda, \mathbf{d}_i)$$

- For each λ_j , we compute the gradient point wise

$$\frac{d}{d \lambda_j} \text{error}(\lambda_j, \mathbf{d}_i) = \frac{d}{d \lambda_j} \text{score}(\lambda, \mathbf{d}_i^{\text{best}}) - \text{score}(\lambda, \mathbf{d}_i)$$

- Gradient with respect to λ_j

$$\frac{d}{d \lambda_j} \text{error}(\lambda_j, \mathbf{d}_i) = \frac{d}{d \lambda_j} \sum_{j'} \lambda_{j'} h_{j'}(\mathbf{d}_i^{\text{best}}) - \sum_{j'} \lambda_{j'} h_{j'}(\mathbf{d}_i)$$

- For $\lambda'_{j'} \neq \lambda_j$, the term $\lambda_{j'} h_{j'}(\mathbf{d}_i)$ is constant, so they disappear

$$\frac{d}{d \lambda_j} \text{error}(\lambda_j, \mathbf{d}_i) = \frac{d}{d \lambda_j} \lambda_j h_j(\mathbf{d}_i^{\text{best}}) - \lambda_j h_j(\mathbf{d}_i)$$

- The derivative of a linear function is its factor

$$\frac{d}{d \lambda_j} \text{error}(\lambda_j, \mathbf{d}_i) = h_j(\mathbf{d}_i^{\text{best}}) - h_j(\mathbf{d}_i)$$

\Rightarrow Our model update is $\lambda_j^{\text{new}} = \lambda_j - (h_j(\mathbf{d}_i^{\text{best}}) - h_j(\mathbf{d}_i))$

- Feature values in model best translation
- Feature values in reference translation
- Intuition:
 - promote features whose value is bigger in reference
 - demote features whose value is bigger in model best

Input: set of sentence pairs (\mathbf{e}, \mathbf{f}), set of features

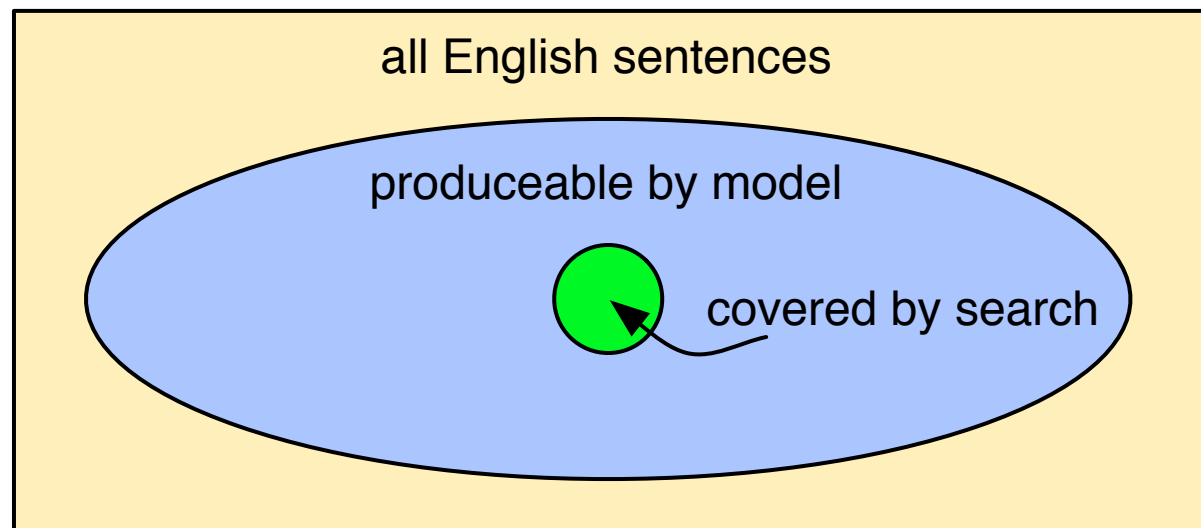
Output: set of weights λ for each feature

```
1:  $\lambda_i = 0$  for all  $i$ 
2: while not converged do
3:   for all foreign sentences  $\mathbf{f}$  do
4:      $\mathbf{e}_{\text{best}} =$  best translation according to model
5:      $\mathbf{e}_{\text{ref}} =$  reference translation
6:     if  $\mathbf{e}_{\text{best}} \neq \mathbf{e}_{\text{ref}}$  then
7:       for all features  $h_i$  do
8:          $\lambda_i += h_i(\mathbf{f}, \mathbf{e}_{\text{ref}}) - h_i(\mathbf{f}, \mathbf{e}_{\text{best}})$ 
9:       end for
10:    end if
11:  end for
12: end while
```

generating the reference


Failure to Generate Reference

- Reference translation may be anywhere in this box




- If producible by model \rightarrow we can compute feature scores
- If not \rightarrow we can not

- Reference translation in tuning set not literal
- Failure even if phrase pairs are extracted from same sentence pair
- Examples



alignment points too distant
→ phrase pair too big to extract



required reordering distance too large
→ exceeds distortion limit of decoder

- BLEU+1
 - add one free n-gram count to statistics → avoids BLEU score of 0
 - however: wrong balance between 1-4 grams, too drastic brevity penalty
- BLEU impact
 - leave all other sentence translations fixed
 - collect n-gram matches and totals from them
 - add n-gram matches and total from current candidate
 - consider impact on overall BLEU score
- Incremental BLEU impact
 - maintain decaying statistics for n-gram matches, total n-grams

$$\text{count}_t = \frac{9}{10} \text{count}_{t-1} + \text{current-count}_t$$

Problems with Max-BLEU Training

- Consider the following Arabic sentence (written left-to-right in Buckwalter romanization) with English glosses:

sd qTEp mn AlkEk AlmmlH " brytzl " Hlqh .
blocked piece of biscuit salted " pretzel " his-throat

- Very literal translation might be

A piece of a salted biscuit, a "pretzel," blocked his throat.

- But reference translation is

A pretzel, a salted biscuit, became lodged in his throat.


- Reference accurate, but major transformations
- Trying to approximate reference translation may lead to bad rules

note: example from Chiang (2012)

mira

Hope and Fear

- Bad: optimize towards **utopian**, away from **n-best**
- Good: optimize towards **hope**, away from **fear**



Hope and Fear Translations

- Hope translation

$$\mathbf{d}^{\text{hope}} = \operatorname{argmax}_{\mathbf{d}} \operatorname{BLEU}(\mathbf{d}) + \operatorname{score}(\mathbf{d})$$

- Finding the fear translation

- Metric difference (should be big)

$$\Delta \operatorname{BLEU}(\mathbf{d}^{\text{hope}}, \mathbf{d}) = \operatorname{BLEU}(\mathbf{d}^{\text{hope}}) - \operatorname{BLEU}(\mathbf{d})$$

- Score difference (should be small or negative)

$$\Delta \operatorname{score}(\lambda, \mathbf{d}^{\text{hope}}, \mathbf{d}) = \operatorname{score}(\lambda, \mathbf{d}^{\text{hope}}) - \operatorname{score}(\lambda, \mathbf{d})$$

- Margin

$$v(\lambda, \mathbf{d}^{\text{hope}}, \mathbf{d}) = \Delta \operatorname{BLEU}(\mathbf{d}^{\text{hope}}, \mathbf{d}) - \Delta \operatorname{score}(\lambda, \mathbf{d}^{\text{hope}}, \mathbf{d})$$

- Fear translation

$$\mathbf{d}^{\text{fear}} = \operatorname{argmax}_{\mathbf{d}} v(\lambda, \mathbf{d}^{\text{hope}}, \mathbf{d})$$

Margin Infused Relaxed Algorithm (MIRA) 34



- Stochastic gradient descent update with learning weight δ_i

$$\lambda_j^{\text{new}} = \lambda_j - \delta_i \left(h_j(\mathbf{d}_i^{\text{fear}}) - h_j(\mathbf{d}_i^{\text{hope}}) \right)$$

- Updates should depend on margin

$$\delta_i = \min \left(C, \frac{\Delta_{\text{BLEU}}(\mathbf{d}_i^{\text{hope}}, \mathbf{d}_i^{\text{fear}}) - \Delta_{\text{score}}(\mathbf{d}_i^{\text{hope}}, \mathbf{d}_i^{\text{fear}})}{\|\Delta h\|^2} \right)$$

- The math behind this is a bit complicated

Different Learning Rates for Features

- For some features, we have a lot of evidence (coarse features)
- Others occur only rarely (sparse features)
- After a while, we do not want to change coarse features too much

⇒ Adaptive Regularization of Weights (AROW)

- record confidence in weights over time
- include this in the learning rate for each feature

- Training is computationally expensive

⇒ Break up training data into batches

- After processing all the batches, average the weights
- Not only a speed-up, also seems to improve quality
- Allows parallel processing, but requires inter-process communication

- Generating hope and fear translations is expensive
 - Sample good/bad by random walk through alignment space
 - use operations as in Gibbs samples
 - vary one translation option choice
 - vary one reordering decision
 - vary one phrase segmentation decision
 - adopt new translation based on relative score
 - Compare current translation against its neighbors
- apply MIRA update if more costly translation has higher BLEU

- MIRA requires translation of each sentence on demand
 - repeated decoding needed
 - computationally very expensive
- Batch MIRA
 - n-best list or search graph (lattice)
 - straightforward parallelization
 - does not seem to harm performance

pro

Scored N-Best List

- Reference translation: he does not go home
- N-best list

Translation	Feature values						BLEU+1
it is not under house	-32.22	-9.93	-19.00	-5.08	-8.22	-5	27.3%
he is not under house	-34.50	-7.40	-16.33	-5.01	-8.15	-5	30.2%
it is not a home	-28.49	-12.74	-19.29	-3.74	-8.42	-5	30.2%
it is not to go home	-32.53	-10.34	-20.87	-4.38	-13.11	-6	31.2%
it is not for house	-31.75	-17.25	-20.43	-4.90	-6.90	-5	27.3%
he is not to go home	-35.79	-10.95	-18.20	-4.85	-13.04	-6	31.2%
he does not home	-32.64	-11.84	-16.98	-3.67	-8.76	-4	36.2%
it is not packing	-32.26	-10.63	-17.65	-5.08	-9.89	-4	21.8%
he is not packing	-34.55	-8.10	-14.98	-5.01	-9.82	-4	24.2%
he is not for home	-36.70	-13.52	-17.09	-6.22	-7.82	-5	32.5%

- Higher quality translation (BLEU+1) should rank higher

Pick 2 Translations at Random

- Reference translation: he does not go home
- N-best list

Translation	Feature values						BLEU+1
it is not under house	-32.22	-9.93	-19.00	-5.08	-8.22	-5	27.3%
he is not under house	-34.50	-7.40	-16.33	-5.01	-8.15	-5	30.2%
it is not a home	-28.49	-12.74	-19.29	-3.74	-8.42	-5	30.2%
it is not to go home	-32.53	-10.34	-20.87	-4.38	-13.11	-6	31.2%
it is not for house	-31.75	-17.25	-20.43	-4.90	-6.90	-5	27.3%
he is not to go home	-35.79	-10.95	-18.20	-4.85	-13.04	-6	31.2%
he does not home	-32.64	-11.84	-16.98	-3.67	-8.76	-4	36.2%
it is not packing	-32.26	-10.63	-17.65	-5.08	-9.89	-4	21.8%
he is not packing	-34.55	-8.10	-14.98	-5.01	-9.82	-4	24.2%
he is not for home	-36.70	-13.52	-17.09	-6.22	-7.82	-5	32.5%

- Higher quality translation (BLEU+1) should rank higher

One is Better than the Other

- Reference translation: he does not go home
- N-best list

Translation	Feature values						BLEU+1
it is not under house	-32.22	-9.93	-19.00	-5.08	-8.22	-5	27.3%
he is not under house	-34.50	-7.40	-16.33	-5.01	-8.15	-5	30.2%
it is not a home	-28.49	-12.74	-19.29	-3.74	-8.42	-5	30.2%
it is not to go home	-32.53	-10.34	-20.87	-4.38	-13.11	-6	31.2%
it is not for house	-31.75	-17.25	-20.43	-4.90	-6.90	-5	27.3%
he is not to go home	-35.79	-10.95	-18.20	-4.85	-13.04	-6	31.2%
he does not home	-32.64	-11.84	-16.98	-3.67	-8.76	-4	36.2%
it is not packing	-32.26	-10.63	-17.65	-5.08	-9.89	-4	21.8%
he is not packing	-34.55	-8.10	-14.98	-5.01	-9.82	-4	24.2%
he is not for home	-36.70	-13.52	-17.09	-6.22	-7.82	-5	32.5%

- Higher quality translation (BLEU+1) should rank higher

Learn from the Pairwise Sample

- Pairwise sample

- $\vec{\text{bad}} = (-31.75, -17.25, -20.43, -4.90, -6.90, -5)$
- $\vec{\text{good}} = (-36.70, -13.52, -17.09, -6.22, -7.82, -5)$

- Learn a classifier

- $\vec{\text{bad}} - \vec{\text{good}} \rightarrow \text{☹}$
- $\vec{\text{good}} - \vec{\text{bad}} \rightarrow \text{☺}$

- Use off the shelf maximum entropy classifier to learn weights for each feature
e.g., MegaM (<http://www.umi.acs.umd.edu/~hal/megam/>)

- Collect samples for each sentence pair in tuning set
- For each sentence, sample 1000-best list for 50 pairwise samples
- Reject samples if difference in BLEU+1 score is too small (≤ 0.05)
- Iterate process
 1. set default weights
 2. generate n-best list
 3. build classifier
 4. adopt classifier weights
 5. go to 2, unless converged

leave one out

Leave One Out Training

- Train initial baseline model
- Force translate the training data:
require decoder to match the reference translation
- Collect statistics over translation rules used
- Leave one out:
do not use translation rules originally collected from current sentence pair
- Related to jackknife
 - 90% of training data used for rule collection
 - 10% to validate rules
 - rotate

Translate Almost All Sentences

- Relaxed leave-one-out
 - allow rules originally collected from current sentence pair
 - very costly → only used, if everything else fails
- Allow single word translations (avoid OOV)
- Larger distortion limit
- Word deletion and insertion (very costly)

Model Re-Estimation

- Generate 100-best list
- Collect fractional counts from derivations

⇒ Much smaller model

⇒ Sometimes better model

Summary

